

Random Forest Regressor-Medical Data Set

Analysis Goals:

Predict the number of days a patient stays in the hospital ("initial_days") utilizing the gathered data.

```
In [8]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
```

```
In [9]: df= pd.read_csv('C:/Users/cynth/OneDrive/Documents/MS Data Analytics/Cleaned_Medical_Da
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 9120 entries, 0 to 9119
```

```
Data columns (total 50 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	9120 non-null	int64
1	Lat	9120 non-null	float64
2	Lng	9120 non-null	float64
3	Population	9120 non-null	int64
4	Children	9120 non-null	int64
5	Age	9120 non-null	int64
6	Income	9120 non-null	float64
7	ReAdmis	9120 non-null	int64
8	VitD_levels	9120 non-null	float64
9	Doc_visits	9120 non-null	int64
10	Full_meals_eaten	9120 non-null	int64
11	vitD_supp	9120 non-null	int64
12	Soft_drink	9120 non-null	int64
13	HighBlood	9120 non-null	int64
14	Stroke	9120 non-null	int64
15	Overweight	9120 non-null	int64
16	Arthritis	9120 non-null	int64
17	Diabetes	9120 non-null	int64
18	Hyperlipidemia	9120 non-null	int64
19	BackPain	9120 non-null	int64
20	Anxiety	9120 non-null	int64
21	Allergic_rhinitis	9120 non-null	int64
22	Reflux_esophagitis	9120 non-null	int64
23	Asthma	9120 non-null	int64
24	Initial_days	9120 non-null	float64
25	TotalCharge	9120 non-null	float64
26	Additional_charges	9120 non-null	float64
27	Item1	9120 non-null	int64
28	Item2	9120 non-null	int64
29	Item3	9120 non-null	int64
30	Item4	9120 non-null	int64
31	Item5	9120 non-null	int64
32	Item6	9120 non-null	int64
33	Item7	9120 non-null	int64
34	Item8	9120 non-null	int64
35	Area_Suburban	9120 non-null	int64
36	Area_Rural	9120 non-null	int64
37	Marital_Married	9120 non-null	int64
38	Marital_Separated	9120 non-null	int64

```

39 Marital_Never Married          9120 non-null    int64
40 Marital_Divorced              9120 non-null    int64
41 Gender_Female                 9120 non-null    int64
42 Gender_Nonbinary              9120 non-null    int64
43 Initial_admin_Elective Admission 9120 non-null    int64
44 Initial_admin_Observation Admission 9120 non-null    int64
45 Complication_risk_High         9120 non-null    int64
46 Complication_risk_Low         9120 non-null    int64
47 Services_Intravenous          9120 non-null    int64
48 Services_CT Scan              9120 non-null    int64
49 Services_MRI                  9120 non-null    int64
dtypes: float64(7), int64(43)
memory usage: 3.5 MB

```

Random Forest Regressor

```

In [10]: from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error as MSE
from sklearn.model_selection import cross_val_score

X = df.drop(['Initial_days'], axis=1)
y = df['Initial_days']

```

```

In [11]: #Split the data into training and test data sets.
X_train, X_test, y_train, y_test= train_test_split(X, y, test_size=0.3, random_state= 4)
X_train.to_csv("D209Task2_X_train.csv")
X_test.to_csv("D209Task2_X_test.csv")
y_train.to_csv("D209Task2_y_train.csv")
y_test.to_csv("D209Task2_y_test.csv")

```

```

In [12]: #instantiate Random Forest Regressor
rf = RandomForestRegressor(max_depth = 6)
#train the model
rf.fit(X_train, y_train)
#predict labels for X_test
cv=cross_val_score(rf,X_train, y_train,scoring='r2', cv=5 )
cv_mean = np.mean(cv)
cv_test=cross_val_score(rf,X_test, y_test,scoring='r2', cv=5 )
cv_test_mean = np.mean(cv_test)
y_pred=rf.predict(X_test)
# Evaluate
mse = MSE(y_test, y_pred)
rmse=(mse**(1/2))
print('Test set MSE of rf: {:.4f}'.format(mse))
print('Test set RMSE of rf: {:.4f}'.format(rmse))
print( 'Cross Val Score R2 for train data: {:.4f}'.format(cv_mean))
print( 'Cross Val Score R2 for test data: {:.4f}'.format(cv_test_mean))
print('rf train score: ', rf.score(X_train, y_train))
print('rf test score', rf.score(X_test, y_test))

```

```


Test set MSE of rf: 4.6822
Test set RMSE of rf: 2.1639
Cross Val Score R2 for train data: 0.9928
Cross Val Score R2 for test data: 0.9926
rf train score: 0.9936514742989756
rf test score 0.9931976076328076

```

```
In [13]: #feature importance for rf tree
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib.pyplot import figure


figure(figsize=(10, 8), dpi=80)
#create a pd.series of feature importance
importances_rf = pd.Series(rf.feature_importances_, index=X.columns)
# sort importances_rf =
sorted_importances_rf = importances_rf.sort_values()
#make horizontal bar plot
sorted_importances_rf.plot(kind='barh', color='lightblue')

plt.show()
```



```
In [14]: sns.scatterplot(x=df.TotalCharge, y=df.Initial_days)
```

```
Out[14]: <AxesSubplot:xlabel='TotalCharge', ylabel='Initial_days'>
```



Data Summary and Implications

1. Accuracy and the mean squared error (MSE) of the prediction model.

The following metrics were calculated for this random forest classifier. As shown below, the MSE is 4.9706, the R2 of both the training and test data was 0.99, and the scores for the random forest model are also 0.99. This all seems to indicate that this model does a good job of predicting initial_days.

```
Test set MSE of rf: 4.9706
Test set RMSE of rf: 2.2295
Cross Val Score R2 for train data: 0.9928
Cross Val Score R2 for test data: 0.9926
rf train score: 0.9935368952029457
rf test score 0.9927787420296704
```

2. Implications of the prediction analysis.

The MSE, RMSE, R2, and model score seem to indicate that the random forest regressor is highly accurate and can be used to help predict the initial days in a visit.

3. Limitations of the data analysis.

One limitation of my analysis is that I did not utilize a grid search to tune the hyperparameters for my random forest regressor and I used the default number of regression trees. Its possible that tuning these hyperparameters and increading the number of trees could make the model better, but since it is already so accurate, it probably wouldn't make much of a difference in terms of helping the client predict initial days.

4. Recommended course of action

As noted in the feature importance plot, one variable , total charge, is almost a perfect fit for initial days. My next steps would be to remove this variable to compare the accuracy and then check in with the client to hear their thoughts. Otherwise, would Since the random forest regressor seems to be highly accurate, I would suggest that the organization can proceed and continue to run the model to predict new data. If the accuracy of the model drops significantly then we could re-train the data and experiment with voting classifiers or decision trees with bootsting.